

Automated Property Value Prediction using MLOps

Stefano Chen
Software Engineer
University of Trieste
Piazzale Europa, 1, 34127 Trieste TS
stefano.chen@studenti.units.it

ABSTRACT

This paper presents an MLOps-based approach for predicting residential property values in California. A local real estate agency relies on manual estimations with an average error rate of 20%. To improve accuracy, we built a complete machine learning pipeline incorporating experiment tracking, model selection, dataset versioning, and web deployment. Among five regression models tested, Random Forest achieved the best results and was further optimized using Optuna. The final model was deployed via a Streamlit web app integrated with MongoDB to log user inputs and predictions. All experiments, datasets, and models were tracked using CometML. The deployed system improves accuracy, transparency, and reproducibility over manual approaches.

CCS Concepts

- **Information system** → **Geographic information systems**
- **Computing methodologies** → **Machine learning**
- **Software and its engineering** → **Software creation and management**

Keywords

MLOps; property valuation; model deployment; CometML; Streamlit; regression models; real estate; Github; Optuna; MongoDB

1. INTRODUCTION

Accurate property valuation is vital in the real estate market, particularly in high-demand regions like California.

A local agency relied on manual expert assessments, which were often time-consuming, inconsistent, and exhibited high error margins ($\approx 20\%$).

We collaborated with the agency to create a machine learning-based solution supported by an MLOps workflow. The system leverages historical housing data and modern tooling to ensure reproducibility, model traceability, and fast deployment. Our pipeline includes dataset versioning, experiment tracking with CometML, hyperparameter tuning using Optuna, and real-time deployment through Streamlit Community Cloud. Additionally, MongoDB was used to log requests and predictions for future auditing and improvement.

The codebase is available at: [GitHub Repository](#)

The live app is accessible at: [Streamlit App](#)

2. RELATED WORK

Several machine learning approaches have been applied to real estate valuation, ranging from simple linear models to advanced ensemble methods. While earlier works emphasized statistical modeling, recent research highlights the importance of operational practices like experiment tracking, versioning, and deployment

pipelines. Despite this, few studies integrate the full MLOps lifecycle in practical housing valuation systems. Our approach fills this gap by combining predictive modeling with modern MLOps tools to deliver a reproducible and scalable solution.

3. METHODOLOGY

3.1 Data Cleaning and Transformation

The dataset used in this project contains detailed records of residential properties across California, featuring variables such as geographical coordinates, house attributes, and proximity to the ocean and major cities. To prepare the data for our regression task, we carried out several cleaning and transformation steps. To address missing values - both within the dataset and in potential user input - we applied a mean imputation. Several numerical features showed right-skewed distribution, and presence of outliers, so we applied log transformation to reduce skewness and enhance the model's sensitivity to variation across a more balanced range. Finally, we standardized all numerical features using a Standard Scaler to improve convergence during model training.

3.2 Feature Engineering

To enhance the model's predictive performance, several composite features were engineered based on domain knowledge and Exploratory Data Analysis (EDA):

- **Rooms_Per_House**: Represents the average number of rooms per house, serving as a proxy for house size.
- **Bedrooms_Ratio**: Indicates the proportion of bedrooms to total rooms, highlighting property layout and potential luxury level.
- **People_Per_House**: Captures neighborhood density and may reflect housing demand in a given area.

These derived features help capture more abstract relationships and interactions within the data. Initial correlation analysis showed that these engineered features had stronger predictive power than some original features.

3.3 Experiment Tracking and Versioning

We employed Github for source control, ensuring collaborative development and reproducibility. CometML was adopted to track experiments, dataset versions, environment configurations, model performance metrics (RMSE, MAE, MAPE and R2) and model registration. This setup enabled efficient comparison and reproducibility.

3.4 Model Selection

Five regression algorithms – Linear Regression, Decision Tree, K-Nearest Neighbors, Support Vector Machine, and Random Forest – were evaluated using 10-fold cross validation with default parameters. RMSE and MAPE were the primary evaluation metrics.

3.5 Hyperparameter Tuning

Random Forest, the top-performing algorithm, underwent hyperparameter optimization with Optuna. Key parameters such as number of estimators, criterion, and minimum samples per leaf were tuned over 100 trials. Every tuning experiment was versioned and tracked in CometML, ensuring traceability and reproducibility. The best configuration was registered for deployment.

3.6 Web Deployment

The final model was deployed using Streamlit web application, offering an intuitive web interface for real-time property value predictions. The app retrieves production-ready models from CometML’s model registry.

3.7 Logging and Monitoring

The web application interacts with a MongoDB instance to log user inputs and corresponding model predictions. This logging mechanism supports continuous monitoring, allowing for future auditing and system improvements.

Additionally, the stored inputs serve as a basis for detecting feature data drift – shifts in input data distribution over time – which is critical for maintaining model performance and triggering retraining when necessary.

4. RESULTS

4.1 Model Selection Outcome

To determine the most suitable model for property value prediction, we conducted a comparative analysis of five widely used regression algorithms: Linear Regression, Decision Tree, K-Nearest Neighbors (KNN), Support Vector Machine (SVM), and Random Forest. Each model was evaluated using the default parameters and a 10-fold cross-validation to ensure robustness and minimize overfitting cases.

The evaluation metrics used are Root Mean Squared Error (RMSE) [Figure 1] and Mean Absolute Percentage Error (MAPE) [Figure 2].




Figure 1. RMSE plot during model selection.




Figure 2. MAPE plot during model selection.

Table 1. Model Selection Metrics

Model	RMSE (avg)	MAPE (avg)
Linear Regression	70700	32%
Decision Tree	65000	22%
KNN	59000	21%
SVM	118400	52%
Random Forest	47000	17%

Linear Regression showed the second poorest performance, likely due to its inability to capture the complex, non-linear relationship within the dataset.

Decision Tree and **KNN** performed moderately well.

SVM struggled with the dimensionality of the dataset, resulting in the poorest performance relative to the other models.

Random Forest achieved the best results across both the metrics, making it the most suitable model for this problem.

4.2 Hyperparameter Tuning Outcome

Using Optuna, we performed automated hyperparameter optimization on the Random Forest model. The search space included:

- n_estimators: 100 – 300
- criterion: [“squared_error”, “absolute_error”, “friedman_mse”, “poisson”]
- min_samples_leaf: 1 – 50

After 100 optimization trials the best configuration found was:

Table 2. Best Hyperparameter

Hyperparameter	Value
Number of Estimators	135
Criterion	“poisson”
Min Samples per Leaf	2

This configuration scored the following metrics value:

Table 3. Best Configuration metrics value

Metric	Value
RMSE	46395.6
MAPE	16.4%
MAE	29509.7
R2	83.6%

The hyperparameter tuning improved marginally the performance of the model, this is most likely due to the low number of optimization trials.

5. LIMITATIONS

While the system is both deployable and modular, it currently lacks an automated retraining mechanism. This limitation is caused by the delayed availability of ground truth labels (i.e., actual property sale prices) for the logged predictions, which are essential for supervised model updates. Consequently, retraining must be performed manually once new labeled data becomes available.

6. CONCLUSION

We developed and deployed a reproducible machine learning workflow to predict houses prices in California.

Random Forest, selected through rigorous model comparison and tuned via Optuna, achieved acceptable performance, allowing the real estate agency to shift resources away from the labor-intensive task of manually estimating house prices.

Future enhancements will focus on integrating a CI/CD pipeline to support automated retraining and further refining hyperparameter optimization strategies.

7. ACKNOWLEDGMENTS

We thank the local real estate agency for providing domain insights and the open-source community for maintaining tools such as Streamlit, Scikit-Learn, and Optuna

8. REFERENCES

- [1] Scikit-learn. Random Forest Regressor Documentation. <https://scikit-learn.org/stable/modules/generated/sklearn.ensemble.RandomForestRegressor.html>.
- [2] Google Cloud. MLOps: Continuous delivery and automation pipelines in machine learning. <https://cloud.google.com/architecture/mlops-continuous-delivery-and-automation-pipelines-in-machine-learning>
- [3] Treveil, M., & Pellegrino, A. (2022). *Practical MLOps: Operationalizing Machine Learning Models*. O'Reilly Media.
- [4] Soriano, F. (2023). California Housing Prices – Extra Features. <https://www.kaggle.com/datasets/fedesoriano/california-housing-prices-data-extra-features>.